




(a) Disaggregated attention and expert execution



Ping-pong micro-batches hide communication behind compute



(b) Reported per-GPU decode gain



M2N communication at 256-KB messages

4.2x

throughput
higher than NCCL

68.2%

median latency
lower than NCCL